

Regularization Techniques

[Code ▾](#)

Sushma Vijayeendra Karjol

Regularization Techniques

Instructor : Dr. Vladimir Shapiro

Introduction

In this module, we will be using the glmnet package to perform the Ridge Regression and the Lasso Regression. To fit the ridge regression models and lasso models the main function of this package, cv.glmnet() or cross-validation glmnet is used. And, this function has quite a different syntax from other model-fitting functions. A college dataset from the ISLR is used for the study. This study aims to know which model fits well to know the predictions of variable - 'Graduation rate' for private and non-private colleges using the rest of the variables as dependents. This is divided into three parts. In the first part, the data set is regularised using the Ridge method, where lambda.min and lambda.1se values are estimated. Then the methods are visually represented and compared. The best fit of the Ridge regression model is set with the train and test dataset. The second part includes the same set of procedures are followed but with LASSO regression. Coefficients of variables are compared for the study to know which one would be the best fit. The third part involves the analysis and discussions.

College Dataset from the library ISLR is assigned to M4_College

[Hide](#)

```
library(ISLR)
attach(College)

M4_College <- College[-1]

# first column has the categorical dataset which is deleted for the study.
```

1) Split the data into a train and test set – refer to the ALY6015_Feature_Selection_R.pdf document for information on how to split a dataset.

The dataset is split into training and testing sets. It is the important part of evaluation data mining models. Typically, a large portion of dataset is allocated for training and smaller for testing.

[Hide](#)

```
library (caret)
```

```
Loading required package: lattice  
Loading required package: ggplot2  
Registered S3 method overwritten by 'data.table':  
  method      from  
  print.data.table
```

[Hide](#)

```
library(lattice)  
library(ggplot2)  
  
# spliiting the dataset for th study  
set.seed(1520)  
trainIndex <- createDataPartition(M4_College$Apps,p=0.7,list=FALSE,times=1)  
M4College_train <- M4_College[ trainIndex,]  
M4College_test <- M4_College[-trainIndex,]
```

The dataset is split into - training and test sets. They differ in size because of the way the samples are being taken.

Train dataset is used to fit the model and the test is conducted to test if the prediction fits the dataset.

[Hide](#)

```
# provides details of the variables  
  
summary(M4College_train)
```

Apps		Accept		Enroll		Top10perc		Top25perc	
Min.	: 81	Min.	: 72	Min.	: 35.0	Min.	: 1.00	Min.	: 9.00
1st Qu.:	776	1st Qu.:	614	1st Qu.:	244.0	1st Qu.:	15.00	1st Qu.:	40.00
Median	: 1558	Median	: 1097	Median	: 430.0	Median	: 24.00	Median	: 55.00
Mean	: 3004	Mean	: 2015	Mean	: 773.4	Mean	: 27.42	Mean	: 55.73
3rd Qu.:	3624	3rd Qu.:	2424	3rd Qu.:	913.0	3rd Qu.:	35.00	3rd Qu.:	69.00
Max.	: 48094	Max.	: 26330	Max.	: 6180.0	Max.	: 96.00	Max.	: 100.00
F.Undergrad		P.Undergrad		Outstate		Room.Board		Books	
Min.	: 139	Min.	: 1.0	Min.	: 2340	Min.	: 1880	Min.	: 96.0
1st Qu.:	979	1st Qu.:	86.0	1st Qu.:	7380	1st Qu.:	3603	1st Qu.:	470.0
Median	: 1707	Median	: 341.0	Median	: 10100	Median	: 4230	Median	: 500.0
Mean	: 3666	Mean	: 869.6	Mean	: 10484	Mean	: 4368	Mean	: 552.7
3rd Qu.:	4296	3rd Qu.:	963.0	3rd Qu.:	12950	3rd Qu.:	5050	3rd Qu.:	600.0
Max.	: 28938	Max.	: 21836.0	Max.	: 20100	Max.	: 8124	Max.	: 2340.0
Personal		PhD		Terminal		S.F.Ratio		perc.alumni	
Min.	: 300	Min.	: 8.00	Min.	: 24.00	Min.	: 2.50	Min.	: 0.00
1st Qu.:	882	1st Qu.:	62.00	1st Qu.:	71.00	1st Qu.:	11.50	1st Qu.:	13.00
Median	: 1200	Median	: 75.00	Median	: 82.00	Median	: 13.60	Median	: 21.00
Mean	: 1362	Mean	: 72.38	Mean	: 79.49	Mean	: 14.04	Mean	: 22.71
3rd Qu.:	1702	3rd Qu.:	85.00	3rd Qu.:	92.00	3rd Qu.:	16.30	3rd Qu.:	31.00
Max.	: 6800	Max.	: 100.00	Max.	: 100.00	Max.	: 39.80	Max.	: 60.00
Expend		Grad.Rate							
Min.	: 3186	Min.	: 10.0						
1st Qu.:	6817	1st Qu.:	53.0						
Median	: 8377	Median	: 66.0						
Mean	: 9612	Mean	: 65.6						
3rd Qu.:	10872	3rd Qu.:	78.0						
Max.	: 45702	Max.	: 118.0						

Hide

```
# dimension of the dataset
dim(M4College_train)
```

```
[1] 545  17
```

Summary determines the descriptive statistics of the all the variables in the dataset.

2) Use the glmnet function to estimate the lambda.min and lambda.1se values. Compare and discuss the values.

The glmnet() function contains alpha argument that provides the model that is fit. If alpha is 0 then a ridge regression model is fit, and if alpha is 1 then a lasso model is fit. First fit Ridge Regression model:

Hide

```
library(glmnet)
```

Hide

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

Hide

```
library(foreach)
#install.packages("Matrix")
library(Matrix)

train.mat = model.matrix(Grad.Rate~., data=M4College_train)[,-1]

lambda_seq <- 10 ^ seq(2,-2,by = -.1)

mod.ridge = cv.glmnet(x = train.mat, y = M4College_train$Grad.Rate, alpha = 0,lambda=
lambda_seq )
mod.ridge
```

Call: cv.glmnet(x = train.mat, y = M4College_train\$Grad.Rate, lambda = lambda_seq, alpha = 0)

Measure: Mean-Squared Error

	Lambda	Index	Measure	SE	Nonzero
min	1.995	18	172	15.77	16
1se	25.119	7	186	17.27	16

Hide

```
cat("\n lambda.best \n")
```

```
lambda.best
```

[Hide](#)

```
lambda.best = mod.ridge$lambda.min  
lambda.best
```

```
[1] 1.995262
```

[Hide](#)

```
cat("\n lambda.lse \n")
```

```
lambda.lse
```

[Hide](#)

```
lambda.lse = mod.ridge$lambda.lse  
lambda.lse
```

```
[1] 25.11886
```

To predict the 'Grad rate variable' train data set is regularized using Ridge method by bringing the co-efficient of all the variables which have a lesser impact on the Graduation rate, are brought to zero. Based on which lambda. min and lambda. lse is calculated. lambda. min is considered as best as it has the least value with 1.995. And, lambda.lse with 25.119 is the largest value of the lambda which has an error of 17.27, usually one standard error away from the minimum. The penalty term lambda regularizes the coefficients to reduce the complexity in the model by shrinking the multi-collinearity.

[Hide](#)

```
test.mat <- model.matrix(Grad.Rate~., data = M4College_test)[,-1]  
y = M4College_test$Grad.Rate  
  
alpha_predict <- predict(mod.ridge, s = mod.ridge$lambda.min , newx = test.mat)
```

since we would compare Ridge model with Lasso in the later part , for consistency we will be using lambda.lse in both the cases.

Hide


```
#glmnet(x, y, alpha = 0.2, weights = wts, nlambda = 20)

fit <- glmnet(train.mat, M4College_train$Grad.Rate , alpha = 0)
```

3) Plot the results from the glmnet function, provide an interpretation. What does this plot tell us?

Hide

```
plot(mod.ridge)
```



In Ridge regression we estimate the coefficients of multiple-regression models where the independent variables are highly correlated. The least square values are unbiased, the variances are large and are far away from the true values. Here, Ridge minimizes the residual sum of squares.

4) Fit a Ridge regression model against the training set and report on the coefficients. Is there anything interesting?

Hide

```
# Rebuilding the model with best lamda value identified
Ridge_best <- glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.best)
Ridge_best
```

```
Call: glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.best)
```

```
   Df %Dev Lambda
1   6 39.14  1.995
```

[Hide](#)

```
# table of co-efficient values for each variable
coef(Ridge_best)
```

```
17 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept) 42.7809929441
Apps         .
Accept       .
Enroll       .
Top10perc    0.0589106769
Top25perc    0.0851650613
F.Undergrad  .
P.Undergrad -0.0004624633
Outstate     0.0013289825
Room.Board   .
Books        .
Personal     -0.0002377431
PhD          .
Terminal     .
S.F.Ratio    .
perc.alumni  0.1432326606
Expend       .
```


[Hide](#)

```
?glmnet()
```

The co-efficient of Apps,Accept,Enroll,Top10per,F.Undergrad,Terminal, Room.Board,Books, Personal,S.F.Ratio approach to zero at a stage stating their lower impact on Grad rate . Further more, the variables Top10perc, Top25perc ,Outstate, perc.alumni have g reater impact on increasing 'Grad rate' in all the Colleges as the coefficients of t hese variables don't tur to zero.


[Hide](#)

```
plot(mod.ridge$glmnet.fit)
```

[Hide](#)

NA
NA

The above plot designates the whole path of variables as they narrow towards zero. As lambda values raise the coefficients approach to zero. Vice-versa when the coefficients are unregularised when the lambda is approaching zero. The ridge regression here has penalized the coefficients such that the least efficient variables are in the estimation to shrink. This indicates that several variables have no impact on the 'graduation rate'. There are two co-efficient that have negative values. The faster the variables are shrinking they have a lesser role to play and can increase the chances of 'Grad rate'.


[Hide](#)

```
# creating matrix for test dataset.
test.mat = model.matrix(Grad.Rate~., data=M4College_test)
```

5) Determine the performance of the fit model against the training set by calculating the root mean square error (RMSE).
 $\text{sqrt}(\text{mean}((\text{actual} - \text{predicted})^2))$

[Hide](#)

```
# RMSE calculaton for train data set

Ridge_best # model created using best lmada value
```

```
Call:  glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.best)

      Df  %Dev Lambda
1    6 39.14  1.995
```

[Hide](#)

```
alpha_predictTrain <- predict(Ridge_best, s = lambda.best , newx = train.mat)

predict <- alpha_predictTrain

cat ("\n Root mean square value of Ridge regression against train dataset \n")
```

Root mean square value of Ridge regression against train dataset

[Hide](#)

```
sqrt(mean((M4College_train$Grad.Rate - predict)^2))
```

```
[1] 13.43388
```

The root mean square value of Ridge Regression method is 13.6257.

6) Determine the performance of the fit model against the test set by calculating the root mean square error (RMSE). Is your model overfit?

[Hide](#)

```
# matrix for train dataset
test.mat = model.matrix(Grad.Rate~., data = M4College_test)[,-1]

y = M4College_test$Grad.Rate

fit <- glmnet(x = test.mat, y = M4College_test$Grad.Rate)

alpha_predictTest <- predict(Ridge_best, s = lambda.best , newx = test.mat)

actualTest <- mod.ridge
predictTest <- alpha_predictTest

cat ("\n Root mean square value of Ridge regression against test dataset \n")
```

Root mean square value of Ridge regression against test dataset

Hide

```
sqrt(mean((M4College_test$Grad.Rate - predictTest)^2))
```

```
[1] 13.37291
```

In comparison with the RMSE of test dataset is 13.52 while train dataset is 13.62. Its difficult to conclude that the model is overfit or underfit as the RMSE values are almost equal

7) Use the cv.glmnet function to estimate the min and lambda.1se values. Compare and discuss the values.

Hide

```
lambda_seqLasso <- 10^seq(2, -2, by = -.1)

mod.lasso = cv.glmnet(train.mat, M4College_train$Grad.Rate, alpha=1 ,lambda = lambda_
seqLasso)
mod.lasso
```

```
Call: cv.glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda_seqLasso,
               alpha = 1)
```

Measure: Mean-Squared Error

	Lambda	Index	Measure	SE	Nonzero
min	0.3162	26	172.1	15.48	11
1se	1.9953	18	185.8	14.70	6

[Hide](#)

```
lambda.bestLasso = mod.lasso$lambda.min
lambda.bestLasso
```

```
[1] 0.3162278
```

[Hide](#)

```
lambda.1se = mod.lasso$lambda.1se
lambda.1se
```


```
[1] 1.995262
```

The minimum Labda is 0.2512 with Standard error of 8.959 which includes 12 variables . Lambda 1se is 1.2589 with Standard error of 7.668 and with only 8 variables.

8) Plot the results from the glmnet function, provide an interpretation. What does this plot tell us?

[Hide](#)

```
plot(mod.lasso)
```



The above graph is the cross validation curve with red dotted line. The plot also has the lower and upper standard deviation curves along with error bars (lambda sequence). There are two vertical dotted lines which indicate λ_{\min} and λ_{1se} . λ_{\min} provides minimum mean of cross validation error. And, λ_{1se} provides cross-validation error one standard deviation away from the λ_{\min} .

9) Fit a LASSO regression model against the training set and report on the coefficients. Do any coefficients reduce to zero? If so, which ones?

Hide

```
lasso_best <- glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.bestLasso)
lasso_best
```

```
Call: glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.bestLasso)
```

```

Df  %Dev Lambda
1 12 45.15 0.2512

```


[Hide](#)

```

#predict(mod.lasso, s = mod.lasso$lambda.1se , newx =)
#predict(mod.lasso, s=lambda.best, type="coefficients")

plot(mod.lasso$glmnet.fit , label = TRUE)

```



The above graph measure the multi-linear complexity of the model. This is graphically explained against train data set. There are several co-efficient values that are blowing-up at the end of the path. Along side coefficient of few variables are suppressed to zero.

[Hide](#)

```
coef(lasso_best)
```

```
17 x 1 sparse Matrix of class "dgCMatrix"
              s0
(Intercept) 40.8651778076
Apps         0.0005533179
Accept       .
Enroll       .
Top10perc    0.0798603795
Top25perc    0.1060521441
F.Undergrad  .
P.Undergrad -0.0014723046
Outstate     0.0013854095
Room.Board   0.0010241267
Books        -0.0012158800
Personal     -0.0015175287
PhD          .
Terminal     -0.0333919968
S.F.Ratio    .
perc.alumni  0.2305534682
Expend       -0.0002690844
```

10) Determine the performance of the fit model against the training set by calculating the root mean square error (RMSE). $\sqrt{\text{mean}((\text{actual} - \text{predicted})^2)}$

[Hide](#)

```
# calculating RMSE

lasso_best <- glmnet(x = train.mat, y = M4College_train$Grad.Rate, lambda = lambda.be
stLasso)

train.mat <- model.matrix(Grad.Rate~., data = M4College_test)[,-1]
y = M4College_test$Grad.Rate

alpha_predictL <- predict(lasso_best, s = mod.lasso$lambda.min , newx = train.mat)

actualTrain <- mod.lasso
predictTrain <- alpha_predictTest

cat("\n Root mean square error of train dataset \n")
```

```
Root mean square error of train dataset
```

[Hide](#)

```
sqrt(mean((M4College_test$Grad.Rate - predictTrain)^2))
```

```
[1] 13.37291
```

11) Determine the performance of the fit model against the test set by calculating the root mean square error (RMSE). Is your model overfit?

[Hide](#)

```
cat("\n Root mean square error of test dataset \n")
```

```
Root mean square error of test dataset
```

[Hide](#)

```
sqrt(mean((M4College_test$Grad.Rate - predictTest)^2))
```

```
[1] 13.37291
```

The RMSE for train dataset and test data set is equivalent with 13.52. Through this it can be said that the model is good fit. However its diffcilut to conclude that it is good fit for any of the independent data sets.

12) Which model performed better and why? Is that what you expected?

The aim of using these two models is to shrink and regularize the co-efficient. This improves the prediction accuracy and decreases the variance. Altogether interpretability can also be improved.

In ridge regression, a penalty is tuning the parameter called lambda was chosen from the cross-validation. It is making the model fit small by decreasing the residual sum of squares also by including a shrinkage penalty. Thus the coefficients that have large values are penalized. And as the lambda is growing, the bias is not changed, however, variance is getting dropped. Ridge is not filtering any variables but selecting all the variables in the final model. The RMSE value for the trained dataset is 13.6257 and the test data set is 13.52526.

On the other hand, the lasso is shrinking the coefficient either approximately to zero or setting the variables directly to zero when the lambda value is big. But lasso does the variable selection. A combination of shrinkage along with a selection of variables makes it a better model than ridge when both the models have. The RMSE value for the trained dataset is 13.52526 and the test data set is 13.52526.

In comparison, since lasso filters and selects few variables, better results were expected. On the other hand, Ridge considers all the variables in reducing the coefficients. However, we observed stable RMSE values for both models.

13) Refer to the ALY6015_Feature_Selection_R.pdf document for how to perform stepwise selection and then fit a model. Did this model perform better or as well as Ridge regression or LASSO? Which method do you prefer and why?

[Hide](#)

```
library(leaps)
library(ISLR)
library(dplyr)
```

Regression subset function or Regsubset() from the leaps package identifies different models with different sizes. We need to only specify the 'nmax' values which inputs the maximum number of variables to be included in the model. Among these best variables are quantified using RSS.

[Hide](#)

```
reg.summary
```

```
Subset selection object
Call: regsubsets.formula(Grad.Rate ~ ., data = M4_College, nvmax = 19)
16 Variables (and intercept)
      Forced in Forced out
Apps          FALSE      FALSE
Accept        FALSE      FALSE
Enroll        FALSE      FALSE
```

Top10perc	FALSE	FALSE
Top25perc	FALSE	FALSE
F.Undergrad	FALSE	FALSE
P.Undergrad	FALSE	FALSE
Outstate	FALSE	FALSE
Room.Board	FALSE	FALSE
Books	FALSE	FALSE
Personal	FALSE	FALSE
PhD	FALSE	FALSE
Terminal	FALSE	FALSE
S.F.Ratio	FALSE	FALSE
perc.alumni	FALSE	FALSE
Expend	FALSE	FALSE

1 subsets of each size up to 16

Selection Algorithm: exhaustive

		Apps	Accept	Enroll	Top10perc	Top25perc	F.Undergrad	P.Undergrad	Outstate
1	(1)	" "	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "	" "
12	(1)	" "	" "	" "	" "	" "	" "	" "	" "
13	(1)	" "	" "	" "	" "	" "	" "	" "	" "
14	(1)	" "	" "	" "	" "	" "	" "	" "	" "
15	(1)	" "	" "	" "	" "	" "	" "	" "	" "
16	(1)	" "	" "	" "	" "	" "	" "	" "	" "

		Room.Board	Books	Personal	PhD	Terminal	S.F.Ratio	perc.alumni	Expend
1	(1)	" "	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "	" "
12	(1)	" "	" "	" "	" "	" "	" "	" "	" "
13	(1)	" "	" "	" "	" "	" "	" "	" "	" "
14	(1)	" "	" "	" "	" "	" "	" "	" "	" "

15	(1)	"*	"*	"*	"*	"	"	"*	"*
16	(1)	"*	"*	"*	"*	"*	"*	"*	"*

The summary function sets the variables that provides the best models. In the above results the good variables are represented by the asterisk. From the above we can conclude that a good model can be created using one variable to max of sixteen variables. Among them Outstate, Top25perc ,perc.alumni , Apps are some of the important variables in an ordinal form.

Hide

```
#fit_subset <- lm(Grad.Rate ~ Outstate + Top25perc + P.Undergrad + perc.alumni + Room
.Board + + Apps)

fit_subset <- step(lm(Grad.Rate ~. , data = M4College_train), direction = 'both', st
eps = 16 , trace = FALSE)

Subset_lm_Train <- predict.lm(fit_subset, newdata = M4College_train)

library(mltools)
SubRMSE_train<- mltools::rmse(preds = Subset_lm_Train,
      actuals = M4College_train$Grad.Rate,
      weights = 1,
      na.rm = FALSE)
cat("\n Root mean square error of best fit of train dataset")
```

Root mean square error of best fit of train dataset

Hide

SubRMSE_train

[1] 12.71618

Hide

```
Subset_lm_Test <- predict.lm(fit_subset, newdata = M4College_test)

cat("\n Root mean square error of best fit of test dataset")
```

Root mean square error of best fit of test dataset

Hide

```
SubRMSE_test <- mltools::rmse(preds = Subset_lm_Test,
  actuals = M4College_test$Grad.Rate,
  weights = 1,
  na.rm    = FALSE)
```

```
SubRMSE_test
```

```
[1] 12.76898
```

[Hide](#)

In a comparison of all the three methods used, Lasso and subset used a lesser number of variables to form a good model. The outState variable is not given a good amount of importance in Ridge and lasso but the Subset regression mentions that a good model can be created using only Outstate Variable with 'Grad Rate'. However, practically it may not be correct to predict and improve business based only on one variable. On the other hand, Lasso gives more importance to 'Pct. new students from top 25% of H.S. class' and 'Pct. alumni who donate'. These are the area of concentration that all the institutions need improvement to increase the Graduation rate.

Conclusion

A variety of models fit using various methods such as: Ridge, Lasso and best subsets. We found that the performance from best to worst was:

1. Linear model : Train - 12.71618 Test - 12.76898
2. Lasso model : Train - 13.37291 Test - 13.37291
3. Ridge model : Train - 13.43388 Test - 13.37291

From the above results, we can conclude that Linear has the least mean square and results in the best model inclusive of all the 16 variables that impact the Graduation rate of the education institution. Thus to improve the Graduation rate, the institution can concentrate on all the areas.

Reference

1. Best Subsets Regression Essentials in R Kassambara & Sfd <http://www.sthda.com/english/articles/37-model-selection-essentials-in-r/155-best-subsets-regression-essentials-in-r/>
(<http://www.sthda.com/english/articles/37-model-selection-essentials-in-r/155-best-subsets-regression-essentials-in-r/>)

regression-essentials-in-r/)

2. Deepika Singh Singh <https://www.pluralsight.com/guides/linear-lasso-and-ridge-regression-with-r>
(<https://www.pluralsight.com/guides/linear-lasso-and-ridge-regression-with-r>)
3. Ridge and Lasso Regression Models <http://wavedatalab.github.io/machinelearningwithr/post4.html>
(<http://wavedatalab.github.io/machinelearningwithr/post4.html>)